

Classical Algorithms Implementation

Anagha H.C.

Sachin Prasanna

Abhayjit Singh

October 2023

1 Denoising Diffusion Probabilistic Model

Diffusion Probabilistic Models, often abbreviated as DDPMs, are a noteworthy category within the realm of generative models. They do so by iteratively introducing controlled noise into an initial input signal. The underlying concept is to acquire a deep understanding of the noise removal process, enabling the generation of entirely fresh and coherent data samples.

Main File:

```
1
2 import torch
3 import logging
4 import os
5 import torch.nn as nn
6 from tqdm import tqdm
7 from torch import optim
8 from utils import *
9 from modules import UNet
10 from matplotlib import pyplot as plt
11 import logging
12 from torch.utils.tensorboard import SummaryWriter
13
14 logging.basicConfig(format="%(asctime)s-%(levelname)s: %(message)s", level=
    logging.INFO, datefmt="%I:%M:%S")
15
16
17 class Diffusion:
18     def __init__(self, ns=10000, strt=1e-5, end=0.02, imsiz=256, device="cuda"):
19         :
20         self.ns = ns
21         self.strt = strt
22         self.end = end
23         self.imsiz = imsiz
24         self.device = device
25
26         self.beta = self.scheduler_noise().to(device)
27         self.alpha = 1. - self.beta
28         self.ahat = torch.cumprod(self.alpha, dim=0)
29
30     def sample_timesteps(self, n):
31         return torch.randint(low=1, high=self.ns, size=(n,))
32
33     def scheduler_noise(self):
34         return torch.linspace(self.strt, self.end, self.ns)
35
36     def noise_images(self, x, t):
37         alp1 = torch.sqrt(self.ahat[t])[:, None, None, None]
38         alp2 = torch.sqrt(1 - self.ahat[t])[:, None, None, None]
39         = torch.randn_like(x)
40         return alp1 * x + alp2* ,
41
42     def sample(self, model, n):
```

```

42     logging.info(f"Sampling_{n}_new_images....")
43     model.eval()
44     with torch.no_grad():
45         x = torch.randn((n, 3, self.imsiz, self.imsiz)).to(self.device)
46         for i in tqdm(reversed(range(1, self.ns)), position=0):
47             t = (torch.ones(n) * i).long().to(self.device)
48             pn = model(x, t)
49             alpha = self.alpha[t][:, None, None, None]
50             ahat = self.ahat[t][:, None, None, None]
51             beta = self.beta[t][:, None, None, None]
52             if i > 1:
53                 noise = torch.randn_like(x)
54             else:
55                 noise = torch.zeros_like(x)
56             x = 1 / torch.sqrt(alpha) * (x - ((1 - alpha) / (torch.sqrt(1 -
57                 ahat)))) * pn + torch.sqrt(beta) * noise
58         model.train()
59         x = (x.clamp(-1, 1) + 1) / 2
60         x = (x * 255).type(torch.uint8)
61         return x
62
63 def train(args):
64     setup_logging(args.run_name)
65     device = args.device
66     dataloader = get_data(args)
67     model = UNet().to(device)
68     optimizer = optim.AdamW(model.parameters(), lr=args.lr)
69     mse = nn.MSELoss()
70     diffusion = Diffusion(imsiz=args.image_size, device=device)
71     logger = SummaryWriter(os.path.join("runs", args.run_name))
72     l = len(dataloader)
73
74     for epoch in range(args.epochs):
75         logging.info(f"epoch:_{epoch}:")
76         pbar = tqdm(dataloader)
77         for i, (images, _) in enumerate(pbar):
78             images = images.to(device)
79             t = diffusion.sample_timesteps(images.shape[0]).to(device)
80             x_t, noise = diffusion.noise_images(images, t)
81             pn = model(x_t, t)
82             loss = mse(noise, pn)
83
84             optimizer.zero_grad()
85             loss.backward()
86             optimizer.step()
87
88             pbar.set_postfix(MSE=loss.item())
89             logger.add_scalar("MSE", loss.item(), global_step=epoch * l + i)
90
91             sampled_images = diffusion.sample(model, n=images.shape[0])
92             save_images(sampled_images, os.path.join("results", args.run_name, f"{
93                 epoch}.jpg"))
94             torch.save(model.state_dict(), os.path.join("models", args.run_name, f"
95                 ckpt.pt"))
96
97 def launch():
98     import argparse
99     parser = argparse.ArgumentParser()
100    args = parser.parse_args()
101    args.run_name = "DDPM_Unconditional"
102    args.epochs = 500

```

```

102     args.batch_size = 12
103     args.image_size = 64
104     args.dataset_path = r"C:\Users\it303project\datasets\landscape_img_folder"
105     args.device = "cuda"
106     args.lr = 3e-4
107     train(args)
108
109
110 if __name__ == '__main__':
111     launch()
112     device = "cuda"
113     model = UNet().to(device)
114     ckpt = torch.load("./working/orig/ckpt.pt")
115     model.load_state_dict(ckpt)
116     diffusion = Diffusion(imsiz=64, device=device)
117     x = diffusion.sample(model, 8)
118     print(x.shape)
119     plt.figure(figsize=(32, 32))
120     plt.imshow(torch.cat([
121         torch.cat([i for i in x.cpu()], dim=-1),
122         ], dim=-2).permute(1, 2, 0).cpu())
123     plt.show()

```

UNet File:

```

1
2
3 import torch
4 import torch.nn as nn
5 import torch.nn.functional as F
6
7
8 class SelfAttention(nn.Module):
9     def __init__(self, chs, size):
10         super(SelfAttention, self).__init__()
11         self.chs = chs
12         self.size = size
13         self.mha = nn.MultiheadAttention(chs, 4, batch_first=True)
14         self.ln = nn.LayerNorm([chs])
15         self.ff_self = nn.Sequential(
16             nn.LayerNorm([chs]),
17             nn.Linear(chs, chs),
18             nn.GELU(),
19             nn.Linear(chs, chs),
20         )
21
22     def forward(self, x):
23         x = x.view(-1, self.chs, self.size * self.size).swapaxes(1, 2)
24         x_ln = self.ln(x)
25         atval, _ = self.mha(x_ln, x_ln, x_ln)
26         atval = atval + x
27         atval = self.ff_self(atval) + atval
28         return atval.swapaxes(2, 1).view(-1, self.chs, self.size, self.size)
29
30
31 class DoubleConv(nn.Module):
32     def __init__(self, inch, outch, midch=None, residual=False):
33         super().__init__()
34         self.residual = residual
35         if not midch:
36             midch = outch
37         self.double_conv = nn.Sequential(
38             nn.Conv2d(inch, midch, kernel_size=3, padding=1, bias=False),
39             nn.GroupNorm(1, midch),

```

```

40         nn.GELU(),
41         nn.Conv2d(midch, outch, kernel_size=3, padding=1, bias=False),
42         nn.GroupNorm(1, outch),
43     )
44
45     def forward(self, x):
46         if self.residual:
47             return F.gelu(x + self.double_conv(x))
48         else:
49             return self.double_conv(x)
50
51
52 class Down(nn.Module):
53     def __init__(self, inch, outch, emb_dim=256):
54         super().__init__()
55         self.maxpool_conv = nn.Sequential(
56             nn.MaxPool2d(2),
57             DoubleConv(inch, inch, residual=True),
58             DoubleConv(inch, outch),
59         )
60
61         self.emb_layer = nn.Sequential(
62             nn.SiLU(),
63             nn.Linear(
64                 emb_dim,
65                 outch
66             ),
67         )
68
69     def forward(self, x, t):
70         x = self.maxpool_conv(x)
71         emb = self.emb_layer(t)[: , :, None, None].repeat(1, 1, x.shape[-2], x.
72             shape[-1])
73         return x + emb
74
75 class Up(nn.Module):
76     def __init__(self, inch, outch, emb_dim=256):
77         super().__init__()
78
79         self.up = nn.Upsample(scale_factor=2, mode="bilinear", align_corners=
80             True)
81         self.conv = nn.Sequential(
82             DoubleConv(inch, inch, residual=True),
83             DoubleConv(inch, outch, inch // 2),
84         )
85
86         self.emb_layer = nn.Sequential(
87             nn.SiLU(),
88             nn.Linear(
89                 emb_dim,
90                 outch
91             ),
92         )
93
94     def forward(self, x, skip_x, t):
95         x = self.up(x)
96         x = torch.cat([skip_x, x], dim=1)
97         x = self.conv(x)
98         emb = self.emb_layer(t)[: , :, None, None].repeat(1, 1, x.shape[-2], x.
99             shape[-1])
100         return x + emb

```

```

100
101 class UNet(nn.Module):
102     def __init__(self, cinn=3, coutt=3, dimensiontime=256, device="cuda"):
103         super().__init__()
104         self.device = device
105         self.dimensiontime = dimensiontime
106         self.inc = DoubleConv(cinn, 64)
107         self.dlayer1 = Down(64, 128)
108         self.salayer1 = SelfAttention(128, 32)
109         self.dlayer2 = Down(128, 256)
110         self.salayer2 = SelfAttention(256, 16)
111         self.dlayer3 = Down(256, 256)
112         self.salayer3 = SelfAttention(256, 8)
113
114         self.bottleneck1 = DoubleConv(256, 512)
115         self.bottleneck2 = DoubleConv(512, 512)
116         self.bottleneck3 = DoubleConv(512, 256)
117
118         self.ulayer1 = Up(512, 128)
119         self.salayer4 = SelfAttention(128, 16)
120         self.ulayer2 = Up(256, 64)
121         self.salayer5 = SelfAttention(64, 32)
122         self.ulayer3 = Up(128, 64)
123         self.salayer6 = SelfAttention(64, 64)
124         self.outc = nn.Conv2d(64, coutt, kernel_size=1)
125
126     def pos_encoding(self, t, chs):
127         inv_freq = 1.0 / (
128             10000
129             ** (torch.arange(0, chs, 2, device=self.device).float() / chs)
130         )
131         posa = torch.sin(t.repeat(1, chs // 2) * inv_freq)
132         posb = torch.cos(t.repeat(1, chs // 2) * inv_freq)
133         posen = torch.cat([posa, posb], dim=-1)
134         return posen
135
136     def forward(self, x, t):
137         t = t.unsqueeze(-1).type(torch.float)
138         t = self.pos_encoding(t, self.dimensiontime)
139
140         x1 = self.inc(x)
141         x2 = self.dlayer1(x1, t)
142         x2 = self.salayer1(x2)
143         x3 = self.dlayer2(x2, t)
144         x3 = self.salayer2(x3)
145         x4 = self.dlayer3(x3, t)
146         x4 = self.salayer3(x4)
147
148         x4 = self.bottleneck1(x4)
149         x4 = self.bottleneck2(x4)
150         x4 = self.bottleneck3(x4)
151
152         x = self.ulayer1(x4, x3, t)
153         x = self.salayer4(x)
154         x = self.ulayer2(x, x2, t)
155         x = self.salayer5(x)
156         x = self.ulayer3(x, x1, t)
157         x = self.salayer6(x)
158         output = self.outc(x)
159         return output
160
161 if __name__ == '__main__':
162     net = UNet(device="cpu")

```

```

163     # net = UNet_conditional(num_classes=10, device="cpu")
164     print(sum([p.numel() for p in net.parameters()]))
165     x = torch.randn(3, 3, 64, 64)
166     t = x.new_tensor([500] * x.shape[0]).long()
167     y = x.new_tensor([1] * x.shape[0]).long()
168     print(net(x, t, y).shape)

```

Utils File:

```

1  import os
2  import torch
3  import torchvision
4  from PIL import Image
5  from matplotlib import pyplot as plt
6  from torch.utils.data import DataLoader
7
8
9  def plot_images(images):
10     plt.figure(figsize=(32, 32))
11     plt.imshow(torch.cat([
12         torch.cat([i for i in images.cpu()], dim=-1),
13         ], dim=-2).permute(1, 2, 0).cpu())
14     plt.show()
15
16
17  def save_images(images, path, **kwargs):
18     grid = torchvision.utils.make_grid(images, **kwargs)
19     ndarr = grid.permute(1, 2, 0).to('cpu').numpy()
20     im = Image.fromarray(ndarr)
21     im.save(path)
22
23
24  def get_data(args):
25     transforms = torchvision.transforms.Compose([
26         torchvision.transforms.Resize(80), # args.image_size + 1/4 *args.
27         image_size
28         torchvision.transforms.RandomResizedCrop(args.image_size, scale=(0.8,
29         1.0)),
30         torchvision.transforms.ToTensor(),
31         torchvision.transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
32     ])
33     dataset = torchvision.datasets.ImageFolder(args.dataset_path, transform=
34     transforms)
35     dataloader = DataLoader(dataset, batch_size=args.batch_size, shuffle=True)
36     return dataloader
37
38
39  def setup_logging(run_name):
40     os.makedirs("models", exist_ok=True)
41     os.makedirs("results", exist_ok=True)
42     os.makedirs(os.path.join("models", run_name), exist_ok=True)
43     os.makedirs(os.path.join("results", run_name), exist_ok=True)

```

2 Generative Models

Generative models are a class of artificial intelligence algorithms that focus on creating data rather than making predictions. They learn the underlying patterns and structures within datasets, enabling them to generate new, synthetic data samples. These models have diverse applications, including generating realistic images, natural language text and audio. The ability to generate data with high fidelity and diversity has significant implications for fields like computer vision, natural language processing, and data augmentation. DDPMs also have a similar aim of learning the distribution of a training data sample

and then generating a new sample that closely resembles it.

3 Generative Adversarial Networks

Generative Adversarial Networks (GANs), introduced by Ian Goodfellow and colleagues in 2014, comprise two crucial neural networks: a generator and a discriminator. The generator and discriminator (aka critic). The generator produces a sample, such as an image, from a latent code. Ideally, the distribution of these images should be indistinguishable from the training distribution, while the discriminator's role is to distinguish real data from generated data. They engage in a competitive game, with the generator refining its output to resemble genuine data and the discriminator enhancing its ability to differentiate. Typically, a GAN consists of two networks: GAN training strikes a balance with a dynamic feedback loop; as the generator improves, the discriminator adapts, fostering ongoing competition. GANs excel in producing highly realistic data for computer vision, art, and data augmentation. In a related context, adversarial nets establish a competitive framework, pitting a generative model against a discriminative model. The generative model aims to craft "counterfeit" samples indistinguishable from genuine data, while the discriminative model detects discrepancies. This approach, trainable using methods like backpropagation and dropout algorithms, holds promise for deep generative modeling, aligning with our exploration of Diffusion Models in this literature review.

```
1
2 import torch
3 import torch.nn as nn
4 import torch.optim as optim
5 import torchvision
6 import torchvision.datasets as datasets
7 from torch.utils.data import DataLoader
8 import torchvision.transforms as transforms
9 from torch.utils.tensorboard import SummaryWriter # to print to tensorboard
10
11
12 class Discriminator(nn.Module):
13     def __init__(self, in_features):
14         super().__init__()
15         self.disc = nn.Sequential(
16             nn.Linear(in_features, 128),
17             nn.LeakyReLU(0.01),
18             nn.Linear(128, 1),
19             nn.Sigmoid(),
20         )
21
22     def forward(self, x):
23         return self.disc(x)
24
25
26 class Generator(nn.Module):
27     def __init__(self, z_dim, img_dim):
28         super().__init__()
29         self.gen = nn.Sequential(
30             nn.Linear(z_dim, 256),
31             nn.LeakyReLU(0.01),
32             nn.Linear(256, img_dim),
33             nn.Tanh(), # normalize inputs to [-1, 1] so make outputs [-1, 1]
34         )
35
36     def forward(self, x):
37         return self.gen(x)
38
39
40 # Hyperparameters etc.
41 device = "cuda" if torch.cuda.is_available() else "cpu"
42 lr = 3e-4
```

```

43 z_dim = 64
44 image_dim = 28 * 28 * 1 # 784
45 batch_size = 32
46 num_epochs = 50
47
48 disc = Discriminator(image_dim).to(device)
49 gen = Generator(z_dim, image_dim).to(device)
50 fixed_noise = torch.randn((batch_size, z_dim)).to(device)
51 transforms = transforms.Compose(
52     [
53         transforms.ToTensor(),
54         transforms.Normalize((0.5,), (0.5,)),
55     ]
56 )
57
58 dataset = datasets.MNIST(root="dataset/", transform=transforms, download=True)
59 loader = DataLoader(dataset, batch_size=batch_size, shuffle=True)
60 opt_disc = optim.Adam(disc.parameters(), lr=lr)
61 opt_gen = optim.Adam(gen.parameters(), lr=lr)
62 criterion = nn.BCELoss()
63 writer_fake = SummaryWriter(f"logs/fake")
64 writer_real = SummaryWriter(f"logs/real")
65 step = 0
66
67 for epoch in range(num_epochs):
68     for batch_idx, (real, _) in enumerate(loader):
69         real = real.view(-1, 784).to(device)
70         batch_size = real.shape[0]
71
72         noise = torch.randn(batch_size, z_dim).to(device)
73         fake = gen(noise)
74         disc_real = disc(real).view(-1)
75         lossD_real = criterion(disc_real, torch.ones_like(disc_real))
76         disc_fake = disc(fake).view(-1)
77         lossD_fake = criterion(disc_fake, torch.zeros_like(disc_fake))
78         lossD = (lossD_real + lossD_fake) / 2
79         disc.zero_grad()
80         lossD.backward(retain_graph=True)
81         opt_disc.step()
82
83
84         output = disc(fake).view(-1)
85         lossG = criterion(output, torch.ones_like(output))
86         gen.zero_grad()
87         lossG.backward()
88         opt_gen.step()
89
90         if batch_idx == 0:
91             print(
92                 f"Epoch [{epoch}/{num_epochs}] Batch {batch_idx}/{len(loader)} \
93                 \
94                 \
95                 \
96                 \
97                 \
98                 \
99                 \
100                LossD: {lossD:.4f}, lossG: {lossG:.4f}"
101             )
102
103             with torch.no_grad():
104                 fake = gen(fixed_noise).reshape(-1, 1, 28, 28)
105                 data = real.reshape(-1, 1, 28, 28)
106                 img_grid_fake = torchvision.utils.make_grid(fake, normalize=True)
107                 img_grid_real = torchvision.utils.make_grid(data, normalize=True)
108
109                 writer_fake.add_image(

```



```

103         "Mnist_Fake_Images", img_grid_fake, global_step=step
104     )
105     writer_real.add_image(
106         "Mnist_Real_Images", img_grid_real, global_step=step
107     )
108     step += 1

```

4 Progressively Growing GAN

Progressive Growing is a methodology employed in the training of Generative Adversarial Networks (GANs) to enhance their proficiency in generating high-resolution images. The fundamental concept underpinning Progressive Growing involves the incremental training of a GAN on images of lower resolutions, subsequently advancing to higher resolutions as the training advances. This approach affords the neural network the opportunity to initially grasp fundamental and less intricate image features, which are progressively refined to produce more intricate and detailed visual content.

config.py

```

1 import torch
2 from math import log2
3
4 START_TRAIN_AT_IMG_SIZE = 256
5 DATASET = 'celeba_hq'
6 CHECKPOINT_GEN = "generator.pth"
7 CHECKPOINT_CRITIC = "discriminator.pth"
8 DEVICE = "cuda" if torch.cuda.is_available() else "cpu"
9 SAVE_MODEL = True
10 LOAD_MODEL = False
11 LEARNING_RATE = 1e-3
12 BATCH_SIZES = [32, 32, 32, 16, 16, 16, 16, 8, 4]
13 IMAGE_SIZE = 512
14 CHANNELS_IMG = 3
15 Z_DIM = 256
16 IN_CHANNELS = 256
17 CRITIC_ITERATIONS = 1
18 LAMBDA_GP = 10
19 NUM_STEPS = int(log2(IMAGE_SIZE / 4)) + 1
20
21 PROGRESSIVE_EPOCHS = [50] * len(BATCH_SIZES)
22 FIXED_NOISE = torch.randn(8, Z_DIM, 1, 1).to(DEVICE)
23 NUM_WORKERS = 4

```

progressiveGAN.py

```

1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 from math import log2
5
6 factors = [1, 1, 1, 1, 1/2, 1/4, 1/8, 1/16, 1/32]
7
8
9 class WSConv2d(nn.Module):
10     def __init__(self, input_channel, out_channel, kernel_size=3, stride=1,
11                 padding=1, gain=2):
12         super(WSConv2d, self).__init__()
13         self.conv = nn.Conv2d(input_channel, out_channel, kernel_size, stride,
14                               padding)
15         self.scale = (gain / (input_channel * (kernel_size ** 2))) ** 0.5
16         self.bias = self.conv.bias
17         self.conv.bias = None

```

```

16
17     # conv layer
18     nn.init.normal_(self.conv.weight)
19     nn.init.zeros_(self.bias)
20
21     def forward(self, x):
22         return self.conv(x * self.scale) + self.bias.view(1, self.bias.shape
23             [0], 1, 1)
24
25 class PixelNorm(nn.Module):
26     def __init__(self):
27         super(PixelNorm, self).__init__()
28         self.epsilon = 1e-8
29
30     def forward(self, x):
31         return x / torch.sqrt(torch.mean(x ** 2, dim=1, keepdim=True) + self.
32             epsilon)
33
34 class CNNBlock(nn.Module):
35     def __init__(self, input_channel, out_channel, pixel_norm=True):
36         super(CNNBlock, self).__init__()
37         self.conv1 = WSConv2d(input_channel, out_channel)
38         self.conv2 = WSConv2d(out_channel, out_channel)
39         self.leaky = nn.LeakyReLU(0.2)
40         self.pn = PixelNorm()
41         self.use_pn = pixel_norm
42
43     def forward(self, x):
44         x = self.leaky(self.conv1(x))
45         x = self.pn(x) if self.use_pn else x
46         x = self.leaky(self.conv2(x))
47         x = self.pn(x) if self.use_pn else x
48         return x
49
50
51 class Generator(nn.Module):
52     def __init__(self, z_dim, in_channels, img_channels=3):
53         super(Generator, self).__init__()
54
55         # initial takes 1x1 -> 4x4
56         self.initial = nn.Sequential(
57             PixelNorm(),
58             nn.ConvTranspose2d(z_dim, in_channels, 4, 1, 0),
59             nn.LeakyReLU(0.2),
60             WSConv2d(in_channels, in_channels, kernel_size=3, stride=1, padding
61                 =1),
62             nn.LeakyReLU(0.2),
63             PixelNorm(),
64         )
65
66         self.initial_rgb = WSConv2d(
67             in_channels, img_channels, kernel_size=1, stride=1, padding=0
68         )
69
70         self.prog_blocks, self.rgb_layers = (
71             nn.ModuleList([]),
72             nn.ModuleList([self.initial_rgb]),
73         )
74
75         for i in range(
76             len(factors) - 1
77         ): # -1 to prevent index error because of factors[i+1]

```

```

76         conv_in_c = int(in_channels * factors[i])
77         conv_out_c = int(in_channels * factors[i + 1])
78         self.prog_blocks.append(CNNBlock(conv_in_c, conv_out_c))
79         self.rgb_layers.append(
80             WSConv2d(conv_out_c, img_channels, kernel_size=1, stride=1,
81                     padding=0)
82         )
83
84     def fade_in(self, alpha, upscaled, generated):
85         # alpha should be scalar within [0, 1], and upscale.shape == generated.
86         # shape
87         return torch.tanh(alpha * generated + (1 - alpha) * upscaled)
88
89     def forward(self, x, alpha, steps):
90         out = self.initial(x)
91
92         if steps == 0:
93             return self.initial_rgb(out)
94
95         for step in range(steps):
96             upscaled = F.interpolate(out, scale_factor=2, mode="nearest")
97             out = self.prog_blocks[step](upscaled)
98
99             final_upscaled = self.rgb_layers[steps - 1](upscaled)
100             final_out = self.rgb_layers[steps](out)
101             return self.fade_in(alpha, final_upscaled, final_out)
102
103 class Discriminator(nn.Module):
104     def __init__(self, z_dim, in_channels, img_channels=3):
105         super(Discriminator, self).__init__()
106         self.prog_blocks, self.rgb_layers = nn.ModuleList([]), nn.ModuleList(
107             [])
108         self.leaky = nn.LeakyReLU(0.2)
109
110         for i in range(len(factors) - 1, 0, -1):
111             conv_in = int(in_channels * factors[i])
112             conv_out = int(in_channels * factors[i - 1])
113             self.prog_blocks.append(CNNBlock(conv_in, conv_out, pixel_norm=
114                 False))
115             self.rgb_layers.append(
116                 WSConv2d(img_channels, conv_in, kernel_size=1, stride=1,
117                         padding=0)
118             )
119
120             self.initial_rgb = WSConv2d(
121                 img_channels, in_channels, kernel_size=1, stride=1, padding=0
122             )
123             self.rgb_layers.append(self.initial_rgb)
124             self.avg_pool = nn.AvgPool2d(
125                 kernel_size=2, stride=2
126             ) # down sampling using avg pool
127
128         self.final_block = nn.Sequential(
129             WSConv2d(in_channels + 1, in_channels, kernel_size=3, padding=1),
130             nn.LeakyReLU(0.2),
131             WSConv2d(in_channels, in_channels, kernel_size=4, padding=0, stride
132                 =1),
133             nn.LeakyReLU(0.2),
134             WSConv2d(

```

```

133         in_channels, 1, kernel_size=1, padding=0, stride=1
134     ),
135 )
136
137 def fade_in(self, alpha, downscaled, out):
138     return alpha * out + (1 - alpha) * downscaled
139
140 def minibatch_std(self, x):
141     batch_statistics = (
142         torch.std(x, dim=0).mean().repeat(x.shape[0], 1, x.shape[2], x.
143             shape[3]))
144     return torch.cat([x, batch_statistics], dim=1)
145
146 def forward(self, x, alpha, steps):
147     cur_step = len(self.prog_blocks) - steps
148
149     out = self.leaky(self.rgb_layers[cur_step](x))
150
151     if steps == 0:
152         out = self.minibatch_std(out)
153         return self.final_block(out).view(out.shape[0], -1)
154
155     downscaled = self.leaky(self.rgb_layers[cur_step + 1](self.avg_pool(x))
156         )
157     out = self.avg_pool(self.prog_blocks[cur_step](out))
158     out = self.fade_in(alpha, downscaled, out)
159
160     for step in range(cur_step + 1, len(self.prog_blocks)):
161         out = self.prog_blocks[step](out)
162         out = self.avg_pool(out)
163
164     out = self.minibatch_std(out)
165     return self.final_block(out).view(out.shape[0], -1)
166
167 if __name__ == "__main__":
168     Z_DIM = 100
169     IN_CHANNELS = 256
170     gen = Generator(Z_DIM, IN_CHANNELS, img_channels=3)
171     critic = Discriminator(Z_DIM, IN_CHANNELS, img_channels=3)
172
173     for img_size in [4, 8, 16, 32, 64, 128, 256, 512, 1024]:
174         num_steps = int(log2(img_size / 4))
175         x = torch.randn((1, Z_DIM, 1, 1))
176         z = gen(x, 0.5, steps=num_steps)
177         assert z.shape == (1, 3, img_size, img_size)
178         out = critic(z, alpha=0.5, steps=num_steps)
179         assert out.shape == (1, 1)
180         print(f"Success! At {img_size}")

```

train.py

```

1 import torch
2 import torch.optim as optim
3 import torchvision.datasets as datasets
4 import torchvision.transforms as transforms
5 from torch.utils.data import DataLoader
6 from torch.utils.tensorboard import SummaryWriter
7 from utils import (
8     gradient_penalty,
9     plot_to_tensorboard,
10    save_checkpoint,
11    load_checkpoint,
12    generate_examples

```

```

13 )
14 from progressive_GAN import Discriminator, Generator
15 from math import log2
16 from tqdm import tqdm
17 import config
18
19 torch.backends.cudnn.benchmarks = True
20
21
22 def get_loader(image_size):
23     transform = transforms.Compose(
24         [
25             transforms.Resize((image_size, image_size)),
26             transforms.ToTensor(),
27             transforms.RandomHorizontalFlip(p=0.5),
28             transforms.Normalize(
29                 [0.5 for _ in range(config.CHANNELS_IMG)],
30                 [0.5 for _ in range(config.CHANNELS_IMG)],
31             ),
32         ]
33     )
34     batch_size = config.BATCH_SIZES[int(log2(image_size / 4))]
35     dataset = datasets.ImageFolder(root=config.DATASET, transform=transform)
36     loader = DataLoader(
37         dataset,
38         batch_size=batch_size,
39         shuffle=True,
40         num_workers=config.NUM_WORKERS,
41         pin_memory=True,
42     )
43     return loader, dataset
44
45
46 def train_fn(
47     critic,
48     gen,
49     loader,
50     dataset,
51     step,
52     alpha,
53     opt_critic,
54     opt_gen,
55     tensorboard_step,
56     writer,
57     scaler_gen,
58     scaler_critic,
59 ):
60     loop = tqdm(loader, leave=True)
61     for batch_idx, (real, _) in enumerate(loop):
62         real = real.to(config.DEVICE)
63         cur_batch_size = real.shape[0]
64
65
66         noise = torch.randn(cur_batch_size, config.Z_DIM, 1, 1).to(config.DEVICE)
67
68         with torch.cuda.amp.autocast():
69             fake = gen(noise, alpha, step)
70             critic_real = critic(real, alpha, step)
71             critic_fake = critic(fake.detach(), alpha, step)
72             gp = gradient_penalty(critic, real, fake, alpha, step, device=config.DEVICE)
73             loss_critic = (

```

```

74         -(torch.mean(critic_real) - torch.mean(critic_fake))
75         + config.LAMBDA_GP * gp
76         + (0.001 * torch.mean(critic_real ** 2))
77     )
78
79     opt_critic.zero_grad()
80     scaler_critic.scale(loss_critic).backward()
81     scaler_critic.step(opt_critic)
82     scaler_critic.update()
83
84
85     with torch.cuda.amp.autocast():
86         gen_fake = critic(fake, alpha, step)
87         loss_gen = -torch.mean(gen_fake)
88
89     opt_gen.zero_grad()
90     scaler_gen.scale(loss_gen).backward()
91     scaler_gen.step(opt_gen)
92     scaler_gen.update()
93
94     # Update alpha and ensure less than 1
95     alpha += cur_batch_size / (
96         (config.PROGRESSIVE_EPOCHS[step] * 0.5) * len(dataset)
97     )
98     alpha = min(alpha, 1)
99
100     if batch_idx % 500 == 0:
101         with torch.no_grad():
102             fixed_fakes = gen(config.FIXED_NOISE, alpha, step) * 0.5 + 0.5
103             plot_to_tensorboard(
104                 writer,
105                 loss_critic.item(),
106                 loss_gen.item(),
107                 real.detach(),
108                 fixed_fakes.detach(),
109                 tensorboard_step,
110             )
111             tensorboard_step += 1
112
113     loop.set_postfix(
114         gp=gp.item(),
115         loss_critic=loss_critic.item(),
116     )
117
118     return tensorboard_step, alpha
119
120
121 def main():
122     gen = Generator(
123         config.Z_DIM, config.IN_CHANNELS, img_channels=config.CHANNELS_IMG
124     ).to(config.DEVICE)
125     critic = Discriminator(
126         config.Z_DIM, config.IN_CHANNELS, img_channels=config.CHANNELS_IMG
127     ).to(config.DEVICE)
128
129     # initialize optimizers and scalers for FP16 training
130     opt_gen = optim.Adam(gen.parameters(), lr=config.LEARNING_RATE, betas=(0.0,
131         0.99))
132     opt_critic = optim.Adam(
133         critic.parameters(), lr=config.LEARNING_RATE, betas=(0.0, 0.99)
134     )
135     scaler_critic = torch.cuda.amp.GradScaler()
136     scaler_gen = torch.cuda.amp.GradScaler()

```

```

136
137 # for tensorboard plotting
138 writer = SummaryWriter(f"logs/gan1")
139
140 if config.LOAD_MODEL:
141     load_checkpoint(
142         config.CHECKPOINT_GEN, gen, opt_gen, config.LEARNING_RATE,
143     )
144     load_checkpoint(
145         config.CHECKPOINT_CRITIC, critic, opt_critic, config.LEARNING_RATE,
146     )
147
148 gen.train()
149 critic.train()
150
151 tensorboard_step = 0
152
153 step = int(log2(config.START_TRAIN_AT_IMG_SIZE / 4))
154 for num_epochs in config.PROGRESSIVE_EPOCHS[step:]:
155     alpha = 1e-5
156     loader, dataset = get_loader(4 * 2 ** step)
157     print(f"Current_image_size: {4*2**step}")
158
159     for epoch in range(num_epochs):
160         print(f"Epoch [{epoch+1}/{num_epochs}]")
161         tensorboard_step, alpha = train_fn(
162             critic,
163             gen,
164             loader,
165             dataset,
166             step,
167             alpha,
168             opt_critic,
169             opt_gen,
170             tensorboard_step,
171             writer,
172             scaler_gen,
173             scaler_critic,
174         )
175
176         if config.SAVE_MODEL:
177             save_checkpoint(gen, opt_gen, filename=config.CHECKPOINT_GEN)
178             save_checkpoint(critic, opt_critic, filename=config.CHECKPOINT_CRITIC)
179
180     step += 1
181
182
183 if __name__ == "__main__":
184     main()

```

utils.py

```

1
2 import torch
3 import random
4 import numpy as np
5 import os
6 import torchvision
7 import config
8 from torchvision.utils import save_image
9 from scipy.stats import truncnorm
10
11

```

```

12
13 def plot_to_tensorboard(
14     writer, loss_critic, loss_gen, real, fake, tensorboard_step
15 ):
16     writer.add_scalar("Loss_Critic", loss_critic, global_step=tensorboard_step)
17
18     with torch.no_grad():
19         img_grid_real = torchvision.utils.make_grid(real[:8], normalize=
20             True)
21         img_grid_fake = torchvision.utils.make_grid(fake[:8], normalize=True)
22         writer.add_image("Real", img_grid_real, global_step=tensorboard_step)
23         writer.add_image("Fake", img_grid_fake, global_step=tensorboard_step)
24
25 def gradient_penalty(critic, real, fake, alpha, train_step, device="cpu"):
26     BATCH_SIZE, C, H, W = real.shape
27     beta = torch.rand((BATCH_SIZE, 1, 1, 1)).repeat(1, C, H, W).to(device)
28     interpolated_images = real * beta + fake.detach() * (1 - beta)
29     interpolated_images.requires_grad_(True)
30
31     # Calculate critic scores
32     mixed_scores = critic(interpolated_images, alpha, train_step)
33
34
35     gradient = torch.autograd.grad(
36         inputs=interpolated_images,
37         outputs=mixed_scores,
38         grad_outputs=torch.ones_like(mixed_scores),
39         create_graph=True,
40         retain_graph=True,
41     )[0]
42     gradient = gradient.view(gradient.shape[0], -1)
43     gradient_norm = gradient.norm(2, dim=1)
44     gradient_penalty = torch.mean((gradient_norm - 1) ** 2)
45     return gradient_penalty
46
47
48 def save_checkpoint(model, optimizer, filename="my_checkpoint.pth.tar"):
49     print("=>Saving checkpoint")
50     checkpoint = {
51         "state_dict": model.state_dict(),
52         "optimizer": optimizer.state_dict(),
53     }
54     torch.save(checkpoint, filename)
55
56
57 def load_checkpoint(checkpoint_file, model, optimizer, lr):
58     print("=>Loading checkpoint")
59     checkpoint = torch.load(checkpoint_file, map_location="cuda")
60     model.load_state_dict(checkpoint["state_dict"])
61     optimizer.load_state_dict(checkpoint["optimizer"])
62
63
64     for param_group in optimizer.param_groups:
65         param_group["lr"] = lr
66
67 def generate_examples(gen, steps, truncation=0.7, n=100):
68     gen.eval()
69     alpha = 1.0
70     for i in range(n):
71         with torch.no_grad():
72             noise = torch.tensor(truncnorm.rvs(-truncation, truncation, size
73                 =(1, config.Z_DIM, 1, 1)), device=config.DEVICE, dtype=torch.

```



```

73         float32)
74         img = gen(noise, alpha, steps)
75         save_image(img*0.5+0.5, f"saved_examples/img_{i}.png")
76     gen.train()

```

5 Individual Contribution

Student Name	Paper Implemented
Sachin Prasanna	Denoising Diffusion Probabilistic Models
Anagha H C	Progressively Growing GANS
Abhayjit Singh Gulati	Generative Adversarial Nets

References

1. <https://arxiv.org/pdf/2006.11239.pdf> - Denoising Diffusion Probabilistic Models
2. <https://arxiv.org/pdf/2102.09672.pdf> - Improved Denoising Diffusion Probabilistic Models
3. <https://arxiv.org/pdf/2105.05233.pdf> - Diffusion Models Beat GANs on Image Synthesis
4. <https://arxiv.org/pdf/1406.2661.pdf> - Generative Adversarial Nets
5. <https://arxiv.org/pdf/1710.10196.pdf> - Progressive Growing of GANS for Improved Quality, Stability and Variation